Waveform-to-waveform generation

Learned / fine-tuned · Frozen · Fixed operation

(a) Speech → environmental audio

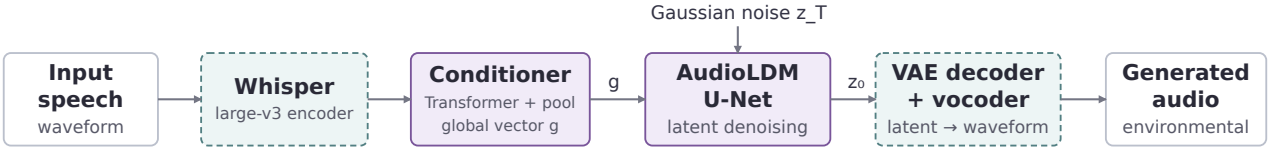

The source speech is encoded directly; no transcript is generated along this path.

(b) Environmental audio → speech

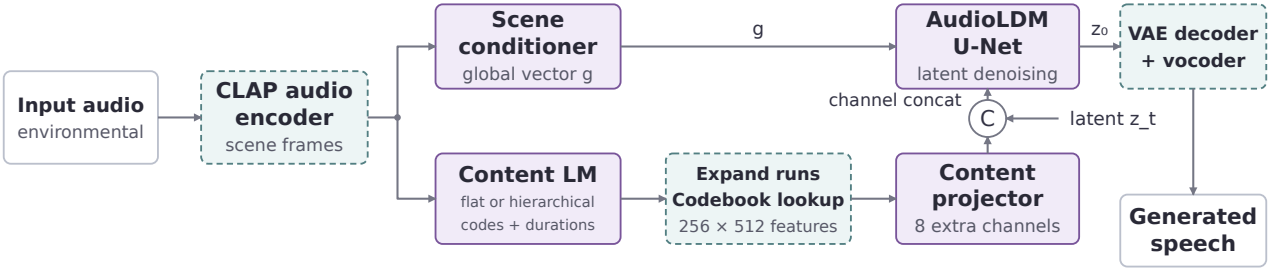

Acoustic content tokens are not words. Target speech and text are used during training, not at inference.